

SHIVAM OMER

Backend Engineer | Python, Django, FastAPI | Microservices & Distributed Systems

Delhi, India | +91-8303669396 | shivamrajomar@gmail.com

shivam-portfolio-gold-omega.vercel.app | linkedin.com/in/programmer-shivam | github.com/shane-Coder

PROFESSIONAL SUMMARY

Backend Engineer with 2.5+ years of experience designing scalable microservices, event-driven architectures, and asynchronous processing pipelines in Python. Skilled in Django, FastAPI, Apache Kafka, Celery, PostgreSQL, and MongoDB, delivering RESTful APIs and high-throughput data pipelines on AWS and Azure. Reduced task execution latency by 80% and improved processing throughput 4x on distributed analytical workloads. Hands-on with Docker, Kubernetes, infrastructure as code, and CI/CD automation on Linux.

TECHNICAL SKILLS

Languages: Python, Golang, Java, C/C++, SQL, Bash/Shell Scripting

Backend & APIs: Django, Django REST Framework, FastAPI, RESTful APIs, Microservices, Celery, SQLAlchemy, JWT Authentication, Asynchronous Programming, System Design

Databases & Messaging: PostgreSQL, MySQL, MongoDB, Redis, Apache Kafka, Data Modeling, Query Optimization

Cloud & DevOps: AWS, Microsoft Azure, Docker, Docker Compose, Kubernetes, Terraform, Ansible, Helm, Jenkins, GitHub Actions, CI/CD Pipelines, Fly.io, Linux

Data & Tools: PySpark, Delta Lake, Azure Databricks, ETL Pipelines, ELK Stack, Git, PyTest, Postman, Agile/Scrum, AI-Assisted Development (Cursor, GitHub Copilot)

PROFESSIONAL EXPERIENCE

Software Development Engineer

Jan 2026 – Present

Move37AI

Delhi, India

- Developed event-driven backend services using Python, Django, Kafka, and MongoDB, processing 10,000+ asynchronous events daily and delivering real-time user notifications in under 5 seconds.
- Designed and delivered 20+ RESTful microservice endpoints for import-export and accounts-receivable platforms, improving transaction orchestration and supporting 3x growth in concurrent internal users.
- Implemented Kafka producer-consumer pipelines with exponential-backoff retry handling, recovering over 95% of transient failures automatically and reducing manual reprocessing tickets.
- Containerized 6+ backend applications with Docker and streamlined Linux deployment workflows, cutting environment setup from ~2 hours to under 15 minutes.
- Accelerated delivery by roughly 30% per feature cycle through AI-assisted engineering workflows (Cursor, GitHub Copilot) and internal automation tooling.

Backend Developer

Mar 2024 – Dec 2025

Solytics Partners

Pune, India

- Engineered WarpDrive, a Python-based distributed task execution framework that reduced task execution latency by nearly 80% across distributed notebook workflows.
- Re-architected asynchronous execution services using Django and Golang, accelerating processing throughput by 4x for large-scale analytical workloads.
- Automated CI/CD pipelines with Jenkins, Groovy, Docker, and Kubernetes, decreasing manual deployment effort by over 90% and cutting release cycles from days to under 1 hour.
- Provisioned cloud-native infrastructure for 10+ services using Terraform, Helm, and AWS, improving deployment consistency and reducing configuration drift incidents.
- Built scalable backend APIs and orchestration layers supporting 100+ concurrent execution pipelines with fault-tolerant retry and microservice coordination.

PROJECTS

PulseCheck — Cron Job Monitoring Service | FastAPI, PostgreSQL, Celery, Redis, Docker, Fly.io · Live Demo · GitHub

- Live production service that detects silently failing cron jobs via dead-man-switch pings, firing email alerts within 60 seconds of a missed check-in, deployed across 5 containerized services (web, worker, scheduler, PostgreSQL, Redis).
- Built 13 REST endpoints with JWT auth over HttpOnly cookies, bcrypt hashing, SQLAlchemy models, and a Celery Beat scheduler, with GitHub Actions CI on every push.

Log Management Microservice | Django, Apache Kafka, MongoDB, Docker · GitHub

- Distributed log ingestion platform processing 10,000+ asynchronous log events with Kafka-based streaming, sub-second real-time alerting, and analytics-ready MongoDB storage across decoupled producer-consumer services.

Azure Databricks Delta Lake Pipeline | Azure Databricks, PySpark, Delta Lake, SQL · GitHub

- Cloud-native ingestion pipeline using Delta Lake ACID tables and micro-batch processing, cutting log transformation runtime by ~60% versus full batch reprocessing through incremental upserts and versioned time-travel queries.

EDUCATION

Bachelor of Technology (B.Tech), Computer Science and Engineering

Aug 2019 – Jun 2023

Dr. A.P.J. Abdul Kalam Technical University

Kanpur, India

ACHIEVEMENTS

Solved 700+ data structures and algorithms problems on LeetCode (Easy, Medium, Hard) and 350+ on GeeksforGeeks.